



Proceedings of the VLDB Endowment

Volume 17, No. 6 – February 2024

Editors in Chief:

Meihui Zhang and Cyrus Shahabi

Associate Editors:

Alkis Polyzotis, Amol Deshpande, Angela Bonifati, Ashraf Aboulnaga, Ashwin Machanavajjhala,
Beng Chin Ooi, Boris Glavic, Ce Zhang, Divy Agrawal, Eric Lo, Fatma Ozcan, Guoliang Li,
Jeffrey Xu Yu, Jian Pei, Jianliang Xu, Johannes Gehrke, K. Selçuk Candan, Kyuseok Shim, Li Xiong,
Magdalena Balazinska, Matthias Boehm, Melanie Herschel, Michael Böhlen,
Nikos Mamoulis, Pinar Tozun, Rachel Pottinger, Sharad Mehrotra, Surajit Chaudhuri, Tamer Özsu,
Tien Tuan Anh Dinh, Walid Aref, Wei Wang, Xiaokui Xiao, Yanyan Shen, Yongxin Tong, Zi Huang

Publication Editors:

Ju Fan, Yang Cao, Xiaoou Ding

PVLDB – Proceedings of the VLDB Endowment

Volume 17, No. 6, February 2024.

All papers published in this issue will be presented at the 50th International Conference on Very Large Data Bases, Guangzhou, China, 2024.

Copyright 2024 VLDB Endowment

This work is licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-nd/4.0/>. For any use beyond those covered by this license, obtain permission by emailing info@vldb.org.

Volume 17, Number 6, February 2024

Pages i – vii and 1159 - 1527

ISSN 2150-8097

Available at: <http://www.pvldb.org> and <https://dl.acm.org/journal/pvldb>

TABLE OF CONTENTS

Front Matter

Copyright Notice	i
Table of Contents	ii
PVLDB Organization and Review Board – Vol. 17	iv

Research Papers

MisDetect: Iterative Mislabeled Detection using Early Loss	1159
<i>Yuhao Deng, Chengliang Chai, Lei Cao, Nan Tang, Jiayi Wang, Ju Fan, Ye Yuan, Guoren Wang</i>	
Capturing More Associations by Referencing External Graphs	1173
<i>Wenfei Fan, Muyang Liu, Shuhao Liu, Chao Tian</i>	
QTCS: Efficient Query-Centered Temporal Community Search.....	1187
<i>Longlong Lin, Pingpeng Yuan, Rong-Hua Li, Chunxue Zhu, Hongchao Qin, Hai Jin, Tao Jia</i>	
DPSUR: Accelerating Differentially Private Stochastic Gradient Descent Using Selective Update and Release	1200
<i>Jie Fu, Qingqing Ye, Haibo Hu, Zhili Chen, Lulu Wang, Kuncan Wang, Xun Ran</i>	
How Can We Train Deep Learning Models Across Clouds and Continents? An Experimental Study	1214
<i>Alexander Erben, Ruben Mayer, Hans-Arno Jacobsen</i>	
Accelerating Sampling and Aggregation Operations in GNN Frameworks with GPU Initiated Direct Storage Accesses.....	1227
<i>Jeongmin Brian Park, Vikram Sharma Mailthody, Zaid Qureshi, Wen-Mei Hwu</i>	
Comprehensive Evaluation of GNN Training Systems: A Data Management Perspective	1241
<i>Hao Yuan, Yajiong Liu, Yanfeng Zhang, Xin Ai, Qiange Wang, Chaoyi Chen, Yu Gu, Ge Yu</i>	
LIION: Fast and High-Resolution Network Kernel Density Visualization.....	1255
<i>Tsz Nam Chan, Rui Zang, Bojian Zhu, Leong Hou U, Dingming Wu, Jianliang Xu</i>	
Performance-Based Pricing of Federated Learning via Auction	1269
<i>Zitao Li, Bolin Ding, Liuyi Yao, Yaliang Li, Xiaokui Xiao, Jingren Zhou</i>	
OEbench: Investigating Open Environment Challenges in Real-World Relational Data Streams.....	1283
<i>Yiqun Diao, Yutong Yang, Qinbin Li, Bingsheng He, Mian Lu</i>	
Influence Maximization via Vertex Countering.....	1297
<i>Jiadong Xie, Zehua Chen, Deming Chu, Fan Zhang, Xuemin Lin, Zhihong Tian</i>	
Optimizing Data Acquisition to Enhance Machine Learning Performance	1310
<i>Tingting Wang, Shixun Huang, Zhifeng Bao, J. Shane Culpepper, Volkan Dedeoglu, Reza Arablouei</i>	
Minimum Strongly Connected Subgraph Collection in Dynamic Graphs	1324
<i>Xin Chen, Jieming Shi, You Peng, Wenqing Lin, Sibao Wang, Wenjie Zhang</i>	
FusionQuery: On-demand Fusion Queries over Multi-source Heterogeneous Data	1337
<i>Junhao Zhu, Yuren Mao, Lu Chen, Congcong Ge, Ziheng Wei, Yunjun Gao</i>	

POLAR: Adaptive and Non-invasive Join Order Selection via Plans of Least Resistance.....	1350
<i>David Justen, Daniel Ritter, Campbell Fraser, Andrew Lamb, Nga Tran, Allison Lee, Thomas Bodner, Mhd Yamen Haddad, Steffen Zeuch, Volker Markl, Matthias Boehm</i>	
DAHA: Accelerating GNN Training with Data and Hardware Aware Execution Planning.....	1364
<i>Zhiyuan Li, Xun Jian, Yue Wang, Yingxia Shao, Lei Chen</i>	
FluidKV: Seamlessly Bridging the Gap between Indexing Performance and Memory-Footprint on Ultra-Fast Storage.....	1377
<i>Ziyi Lu, Qiang Cao, Hong Jiang, Yuxing Chen, Jie Yao, Anqun Pan</i>	
How do Categorical Duplicates Affect ML? A New Benchmark and Empirical Analyses	1391
<i>Vraj Shah, Thomas J Parashos, Arun Kumar</i>	
CGgraph: An Ultra-fast Graph Processing System on Modern Commodity CPU-GPU Co-processor	1405
<i>Pengjie Cui, Haotian Liu, Bo Tang, Ye Yuan</i>	
FCBench: Cross-Domain Benchmarking of Lossless Compression for Floating-point Data	1418
<i>Xinyu Chen, Jiannan Tian, Ian Beaver, Cynthia Freeman, Yan Yan, Jianguo Wang, Dingwen Tao</i>	
PairwiseHist: Fast, Accurate, and Space-Efficient Approximate Query Processing with Data Compression	1432
<i>Aaron Hurst, Daniel E Lucani, Qi Zhang</i>	
MetaStore: Analyzing Deep Learning Meta-Data at Scale.....	1446
<i>Huayi Zhang, Binwei Yan, Lei Cao, Samuel Madden, Elke Rundensteiner</i>	
RTScan: Efficient Scan with Ray Tracing Cores.....	1460
<i>Yangming Lv, Kai Zhang, Ziming Wang, Xiaodong Zhang, Rubao Lee, Zhenying He, Yinan Jing, X. Sean Wang</i>	
FreshGNN: Reducing Memory Access via Stable Historical Embeddings for Graph Neural Network Training	1473
<i>Kezhao Huang, Haitian Jiang, Minjie Wang, Guangxuan Xiao, David Wipf, Xiang Song, Quan Gan, Zengfeng Huang, Jidong Zhai, Zheng Zhang</i>	
Sorting on Byte-Addressable Storage: The Resurgence of Tree Structure	1487
<i>Ying Zheng, Kian-Lee Tan</i>	
Efficient Placement of Decomposable Aggregation Functions for Stream Processing over Large Geo-Distributed Topologies.....	1501
<i>Xenofon Chatziliadis, Eleni Tzirita Zacharatou, Alphan Eracar, Steffen Zeuch, Volker Markl</i>	
AeonG: An Efficient Built-in Temporal Support in Graph Databases	1515
<i>Jiamin Hou, Zhanhao Zhao, Zhouyu Wang, Wei Lu, Guodong Jin, Dong Wen, Xiaoyong Du</i>	

PVLDB ORGANIZATION AND REVIEW BOARD - Vol. 17

Editors in Chief of PVLDB

Meihui Zhang (Beijing Institute of Technology)
Cyrus Shahabi (University of Southern California)

Associate Editors of PVLDB

Alkis Polyzotis (Databricks)
Amol Deshpande (University of Maryland at College Park)
Angela Bonifati (Lyon 1 University)
Ashraf Aboulnaga (Qatar Computing Research Institute, HBKU)
Ashwin Machanavajjhala (Duke)
Beng Chin Ooi (NUS)
Boris Glavic (Illinois Institute of Technology)
Ce Zhang (ETH)
Divy Agrawal (University of California, Santa Barbara)
Eric Lo (Chinese University of Hong Kong)
Fatma Ozcan (Google)
Guoliang Li (Tsinghua University)
Jeffrey Xu Yu (Chinese University of Hong Kong)
Jian Pei (Simon Fraser University)
Jianliang Xu (Hong Kong Baptist University)
Johannes Gehrke (Microsoft)
K. Selçuk Candan (Arizona State University)
Kyuseok Shim (Seoul National University)
Li Xiong (Emory University)
Magdalena Balazinska (UW)
Matthias Boehm (Technische Universität Berlin)
Melanie Herschel (Universität Stuttgart)
Michael Böhlen (University of Zurich)
Nikos Mamoulis (University of Ioannina)
Pinar Tozun (IT University of Copenhagen)
Rachel Pottinger (Univ. of British Columbia)
Sharad Mehrotra (U.C. Irvine)
Surajit Chaudhuri (Microsoft)

Tamer Özsu (University of Waterloo)
Tien Tuan Anh Dinh (Deakin University)
Walid Aref (Purdue University)
Wei Wang (ByteDance)
Xiaokui Xiao (National University of Singapore)
Yanyan Shen (Shanghai Jiao Tong University)
Yongxin Tong (Beihang University)
Zi Huang (University of Queensland)

Publication Editors

Ju Fan (Renmin University of China)
Yang Cao (Hokkaido University)
Xiaou Ding (Harbin Institute of Technology)

PVLDB Managing Editor

Wolfgang Lehner (Dresden University of Technology)

PVLDB Advisory Board

Vanessa Braganholo (Universidade Federal Fluminense)
Sourav S Bhowmick (Nanyang Technological University)
Torsten Grust (University of Tuebingen)
Xin Luna Dong (Facebook)
Fatma Ozcan (Google)
Lei Chen (Hong Kong University of S&T)
Juliana Freire (New York University)
Graham Cormode (University of Warwick)
Divesh Srivastava (AT&T Labs-Research)
Felix Naumann (HPI)
Georgia Koutrika (Athena Research Center)
Jun Yang (Duke University)
Meihui Zhang (Beijing Institute of Technology)
Cyrus Shahabi (University of Southern California)
Nesime Tatbul (Intel Labs and MIT)
Themis Palpanas (Universite Paris Cite)

Review Board

Abolfazl Asudeh (University of Illinois Chicago)
Aditya Parameswaran (University of California, Berkeley)
Ahmed S. Abdelhamid (Purdue University)
Ahmed Eldawy (University of California, Riverside)
Ahmed El-Roby (Carleton University)
Ahmed Mahmood (Google)
Alberto Lerner (University of Fribourg, Switzerland)
Alexander Thomson (Google)
Amr Magdy (University of California Riverside)
Andreas Züfle (Emory University)
Angelos Christos Anadiotis (Oracle)
Anja Gruenheid (Microsoft)
Anthony Tung (National U. of Singapore)
Anton Dignös (Free University of Bozen-Bolzano, Italy)
Arijit Khan (Aalborg University)
Avrilia Floratou (Microsoft)
Baihua Zheng (Singapore Management University)
Bailu Ding (Microsoft Research)
Berthold Reinwald (IBM Research-Almaden)
Bin Yang (East China Normal University)
Bingsheng He (National University of Singapore)
Bolin Ding (Data Analytics and Intelligence Lab, Alibaba Group)
Brandon Haynes (Microsoft Gray Systems Lab)
Chao Zhang (University of Waterloo)
Cheng Long (Nanyang Technological University)
Chengfei Liu (Swinburne University of Technology)
Chengkai Li (The University of Texas at Arlington)
Chengliang Chai (Beijing Institute of Technology)
Chrysanthi Kosyfaki (University of Ioannina)
Chunwei Liu (MIT)
Cong Yan (Microsoft research)
Daisy Zhe Wang (University of Florida)
Dan Kifer (Pennsylvania State Univ., USA)
Dan Lin (Vanderbilt University)
Daniel Kang (UIUC)
Demetrios Zeinalipour-Yazti (University of Cyprus)
Dimitris Papadias (HKUST)
Dong Deng (Rutgers University - New Brunswick)
Dong Wen (University of New South Wales)
Dong Xie (Penn State University)
Dongxiang Zhang (Zhejiang University)
Dumitrel Loghin (National University of Singapore)
Egemen Tanin (University of Melbourne)
El Kindi Rezig (Massachusetts Institute of Technology)
Elena Ferrari (University of Insubria, Varese)
Eser Kandogan (Megagon Labs)
Essam Mansour (Concordia University)
Fan Zhang (Guangzhou University)
Fatemeh Nargesian (University of Rochester)
Fei Chiang (McMaster University)
Feng Zhang (Renmin University of China)
Florin Rusu (UC Merced)
Gabriel Ghinita (Hamad Bin Khalifa University)
Gao Cong (Nanyang Technological University)
George Fakas (Uppsala University)
Haibo Hu (Hong Kong Polytechnic University)
Holger Pirk (Imperial College)
Hong Cheng (Chinese University of Hong Kong)
Hongzhi Wang (Harbin Institute of Technology)
Hua Lu (Roskilde University)
Huanchen Zhang (Tsinghua University)
Huiping Cao (New Mexico State University)
Ibrahim Sabek (MIT)
Ilaria Bartolini (University of Bologna)
Jana Giceva (TU Munich)
Jennie Rogers (Northwestern University)
Jia Zou (Arizona State University)
Jian Lou (Zhejiang University)
Jiangshan Yu (Monash University)
Jianguo Wang (Purdue University)
Jiannan Wang (Simon Fraser University)
Jianqiu Xu (Nanjing University of Aeronautics and Astronautics)
Jianxin Li (Deakin University)
Jieming Shi (The Hong Kong Polytechnic University)
Jin Wang (Megagon Labs)
Jinfei Liu (Zhejiang University)
Johes Bater (Tufts University)
John Liagouris (Boston University)
Jonathan Goldstein (Microsoft)
Ju Fan (Renmin University of China)
Juhua Hu (University of Washington)
Kai Wang (Shanghai Jiao Tong University)
Kangfei Zhao (Beijing Institute of Technology)
Karima Echihabi (Mohammed VI Polytechnic University)
Katja Hose (TU Wien)
Khuzaima Daudjee (University of Waterloo)
Kyoungmin Kim (POSTECH)
Lawrence Benson (HPI, University of Potsdam)
Lei Chen (Hong Kong University of Science and Technology)
Lei Zou (Peking University)
Leong Hou U (University of Macau)
Lin Ma (University of Michigan)
Linyang Chu (McMaster University)
Liyue Fan (UNC Charlotte)
Lu Chen (Zhejiang University)
Luigi Bellomarini (Banca d'Italia)
Madelon Hulsebos (University of Amsterdam)
Manolis Terrovitis (IMIS, Athena RC)
Marco Patella (University of Bologna)
Mario Nascimento (Northeastern University)
Matteo Lissandrini (Aalborg University)
Matthias Renz (University of Kiel)
Michael Hay (Colgate University & Tumult Labs)
mingjie tang (Ant Financial)
Mirek Riedewald (Northeastern University)
Mohamed S. Hassan (Google)
Mohamed Mokbel (University of Minnesota - Twin Cities)
Mohammad Javad Amiri (University of Pennsylvania)
Mostafa Milani (The University of Western Ontario)
Mourad OUZZANI (Qatar Computing Research Institute, HBKU)
Nesime Tatbul (Intel Labs and MIT)
Norman May (SAP SE)
Oliver A Kennedy (University at Buffalo, SUNY)

Panagiotis Bouros (Johannes Gutenberg University Mainz)
 Papotti Paolo (EURECOM)
 Patrick Damme (Technische Universität Berlin)
 Peng Peng (Hunan University)
 Philippe Bonnet (IT Univ Copenhagen, Denmark)
 Pinar Karagoz (METU, Turkey)
 Prashant Pandey (University of Utah)
 Primal Pappachan (Penn State University)
 Qichen Wang (Hong Kong Baptist University)
 Qing Liu (Zhejiang University)
 Qun Chen (Northwestern Polytechnical University)
 Renata Borovica-Gajic (University of Melbourne)
 Rihan Hai (TU Delft)
 Ritesh Ahuja (Oracle Labs)
 Roger Zimmermann (NUS)
 Ronghua Li (Beijing Institute of Technology)
 Sai Wu (Zhejiang Univ)
 Sanjay Krishnan (UChicago)
 Senjuti Basu Roy (NJIT)
 Seokki Lee (University of Cincinnati)
 Shantanu Sharma (New Jersey Institute of Technology)
 Shaofeng Cai (National University of Singapore)
 Shaoxu Song (Tsinghua University)
 Shuai Ma (Beihang University)
 Shuang Hao (Beijing Jiaotong University)
 Sibow Wang (The Chinese University of Hong Kong)
 Stefania Dumbrava (ENSIIE)
 Stefano Paraboschi (Universita' degli Studi di Bergamo)
 Sujaya Maiyya (University of Waterloo)
 Tarique Siddiqui (Microsoft Research)
 Thanaa Ghanem (Metro State University)
 Thang Dinh (VCU)
 Themis Palpanas (Universite Paris Cite)
 Thomas Neumann (TUM)
 Tianhao Wang (University of Virginia)
 Tianzheng Wang (Simon Fraser University)
 Tieying Zhang (Bytedance)
 Tristan Allard (Univ Rennes, CNRS, IRISA)
 Umar Farooq Minhas (Apple)
 Utku Sirin (Harvard University)
 Viktor Leis (Technische Universität München)
 Vincenzo Gulisano (Chalmers University of Technology)
 Vraj Shah (IBM Research)
 Wang-Chien Lee (Pennsylvania State University, USA)
 WEI LU (Renmin University of China)

Wei Wang (Hong Kong University of Science and Technology (Guangzhou))
 Wei-Shinn Ku (Auburn University)
 Wenchao Zhou (Alibaba Group)
 Wendy Hui Wang (Stevens Institute of Technology)
 Xiang Lian (Kent State University)
 Xiang Zhao (National University of Defence Technology)
 Xiangyao Yu (University of Wisconsin-Madison)
 Xiao Hu (Duke University)
 Xiao Hu (University of Waterloo)
 Xiaochun Yang (Northeastern University)
 Xiaofang Zhou (The Hong Kong University of Science and Technology)
 Xiaofei Zhang (University of Memphis)
 Xiaohui Yu (York University)
 Xiaoli Wang (Xiamen University)
 Xin Huang (Hong Kong Baptist University)
 Xin Wang (Tianjin University)
 Xingquan Zhu (Florida Atlantic University)
 Yanfeng Zhang (Northeastern University)
 Yang Cao (Hokkaido University)
 Yannis Chronis (Google)
 Yao Lu (Microsoft Research)
 Ye Yuan (Beijing Institute of Technology)
 Yeye He (Microsoft Research)
 Ying Zhang (University of Technology Sydney)
 Yingxia Shao (BUPT)
 Yu Yang (City University of Hong Kong)
 Yuhao Zhang (University of California, San Diego)
 Yuncheng Wu (National University of Singapore)
 Yunjun Gao (Zhejiang University)
 Yuval Moskovitch (Ben Gurion University)
 Yuxiang Zeng (Beihang University)
 Zhaojing Luo (National University of Singapore)
 Zhengjie Miao (Duke University)
 Zhichao Cao (Arizona State University)
 Zhifeng Bao (RMIT University)
 Zhiwei Zhang (Beijing Institute of Technology)
 Zhongle Xie (Zhejiang University)
 Zhuoyue Zhao (University at Buffalo - SUNY)
 Ziawasch Abedjan (Leibniz Universität Hannover)
 Ziliang Lai (Chinese University of Hong Kong)
 Zimu Zhou (City University of Hong Kong)

LETTER FROM THE EDITORS IN CHIEF

It is our pleasure to present the sixth issue of Volume 17 of PVLDB (Proceedings of the VLDB). PVLDB is dedicated to showcasing original research papers that encompass a wide spectrum of subjects within the area of data and information management. Our coverage spans from fundamental theoretical principles and cutting-edge system architectures to innovative models, techniques, novel applications, and the comprehensive assessment and deployment of large-scale solutions. In our research track, we feature four equally significant categories of papers: (a) regular research, (b) scalable data science (SDS), (c) experiment, analysis & benchmark (EA&B), and (d) vision papers.

The sixth issue of PVLDB's Volume 17 includes 27 papers, spanning the topics of Machine Learning, AI, and Databases; Graph and Network Data; Database Engines; Novel Database Architectures; Data Mining and Analytics; Data Privacy and Security; Information Integration and Data Quality; and User Interfaces. Four topics stood out, the most popular ones in this issue: AI/ML and databases (8 papers), graph and network data (6 papers), database engines (5 papers) and novel database architectures (3 papers).

Out of the 27 papers, six are in the experiment, analysis & benchmark category, and the rest are regular research papers. Twelve papers were accepted after a shepherding round, and the others were accepted after revision.

PVLDB is committed to providing valuable and constructive feedback through a rigorous review process. All submissions undergo meticulous peer review by a team of accomplished Associate Editors and dedicated reviewers. Each paper receives comprehensive evaluation from a minimum of three reviewers, along with the oversight of an Associate Editor. During a three-week discussion phase, reviewers engage in a thorough exchange of perspectives, ultimately converging on a consensus, which is summarized in a meta-review. Some submissions may proceed to a revision phase, affording authors a three-month window to refine their work for subsequent review cycles.

This issue is the result of all the work put in by the authors as well as the great commitment and effort of our associate editors and reviewers as well as our proceedings chairs.

Meihui Zhang and Cyrus Shahabi
Editors-in-Chief of PVLDB Vol. 17
Program Chairs for VLDB 2024